

PROJECT 4: INTRODUCTION TO DEEP LEARNING

Avinash Kommineni, 50248877

December 6, 2017

1 Introduction

The assignment include reading the data, building an optimum model and optimising the hyperparameters. Tensorflow is the library being used and for computational power Google Cloud Platform has been used.

The initialisation and CELEB dataset loading...

Each image is of size 218 X 178.

Initialising some Constants

```
1 usingImages = 4000
2 nClasses = 2
3 shape1 = 178
4 shape2 = 218
5 printEvery = 1
6 batchSize = 100
7 lambdada = 0.01
```

2 Explanations:

Labels

Although there are 202599 images, loading all of them into the physical memory is not recommended and also not possible. So only a portion of these are currently being loaded.

The labels are obtained from text file.

```
1 data = pd.read_csv('files/Anno/list_attr_celeba.txt', ←
    delim_whitespace = True, header=1)
2 df = data['Eyeglasses']
3 df = (df + 1)/2
```

```
4 labels = np.eye(2)[df.values.astype(int)]
```

The above code includes reading the text file and loading all of it into dataframe. It is then cut-down to only one column of 'Eyeglasses'. Since the labels are +1, -1 as True and False, these are modified into one-hot vector.

Images

The images are loaded from the folder as follows...

```
1 images = np.array([np.float32(np.array(Image.open("files/Img/↵
    img_align_celeba/"+str(fname)).resize((shape1, shape2))))/256 ↵
    for fname in df.head(usingImages).index])
2 with open('my.pickle', 'wb') as handle:
3 pickle.dump(images, handle, protocol=pickle.HIGHEST_PROTOCOL)
```

The data loaded is being stored in a pickle 'my.pickle' so that it can be loaded fastly later.

2.1 Convolutional Neural Network architecture

```
1 def convolutionalNeuralNetwork(x):
2     weights = {'conv1':tf.Variable(tf.random_normal([5,5,3,64])),
3               'conv2':tf.Variable(tf.random_normal([5,5,64,128])),
4               'conv3':tf.Variable(tf.random_normal([5,5,128,256])),
5               'conv4':tf.Variable(tf.random_normal([5,5,256,256])),
6               'fullyC1':tf.Variable(tf.random_normal([14*12*256,1024])),
7               'fullyC2':tf.Variable(tf.random_normal([1024,1024])),
8               'out':tf.Variable(tf.random_normal([1024,nClasses]))}
9
10    biases = {'conv1':tf.Variable(tf.random_normal([64])),
11             'conv2':tf.Variable(tf.random_normal([128])),
12             'conv3':tf.Variable(tf.random_normal([256])),
13             'conv4':tf.Variable(tf.random_normal([256])),
14             'fullyC1':tf.Variable(tf.random_normal([1024])),
15             'fullyC2':tf.Variable(tf.random_normal([1024])),
16             'out':tf.Variable(tf.random_normal([nClasses]))}
17
```

```

18 conv1 = tf.nn.relu(tf.nn.conv2d(input=x, filter=weights['conv1']↵
    ],strides=[1,1,1,1],padding='SAME')+ biases['conv1'])
19 conv1 = tf.nn.max_pool(conv1,ksize=[1,2,2,1] ,strides↵
    =[1,2,2,1], padding='SAME')
20 print(conv1)
21
22 conv2 = tf.nn.relu(tf.nn.conv2d(input=conv1, filter=weights['↵
    conv2'],strides=[1,1,1,1],padding='SAME') + biases['conv2'])
23 conv2 = tf.nn.max_pool(conv2,ksize=[1,2,2,1] ,strides↵
    =[1,2,2,1], padding='SAME')
24 print(conv2)
25
26 conv3 = tf.nn.relu(tf.nn.conv2d(input=conv2, filter=weights['↵
    conv3'],strides=[1,1,1,1],padding='SAME') + biases['conv3'])
27 conv3 = tf.nn.max_pool(conv3,ksize=[1,2,2,1] ,strides↵
    =[1,2,2,1], padding='SAME')
28 print(conv3)
29
30 conv4 = tf.nn.relu(tf.nn.conv2d(input=conv3, filter=weights['↵
    conv4'],strides=[1,1,1,1],padding='SAME') + biases['conv4'])
31 conv4 = tf.nn.max_pool(conv4,ksize=[1,2,2,1] ,strides↵
    =[1,2,2,1], padding='SAME')
32 print(conv4)
33
34 conv4 = tf.reshape(conv4,[-1,14*12*256])
35 print(conv4)
36 fcLayer1 = tf.nn.relu(tf.matmul(conv4,weights['fullyC1']) + ↵
    biases['fullyC1'])
37 print(fcLayer1)
38 fcLayer1 = tf.nn.dropout(fcLayer1,keepRate)
39 fcLayer2 = tf.nn.relu(tf.matmul(fcLayer1,weights['fullyC2']) + ↵
    biases['fullyC2'])
40 output = tf.matmul(fcLayer2,weights['out']) + biases['out']
41 print(output)
42
43 regLoss = tf.nn.l2_loss(weights['conv1']) + tf.nn.l2_loss(↵
    weights['conv2']) + tf.nn.l2_loss(weights['conv3']) + tf.nn.↵
    l2_loss(weights['conv3']) + tf.nn.l2_loss(weights['fullyC1'↵
    ]) + tf.nn.l2_loss(weights['fullyC1']) + tf.nn.l2_loss(↵
    weights['out'])
44 return output, regLoss

```

Training

```
1 def trainNetwork():
2     loss = tf.reduce_mean(tf.nn.sigmoid_cross_entropy_with_logits(↵
        logits=tf.transpose(y),labels=tf.transpose(t)) + lambdaa*↵
        regLoss)
3     optimizer = tf.train.AdamOptimizer().minimize(loss)
4     correct = tf.equal(tf.argmax(y),tf.argmax(t))
5     accuracy = tf.reduce_mean(tf.cast(correct,'float'))
6
7     init = tf.global_variables_initializer()
8     sess = tf.Session()
9
10    sess.run(init)
11
12    for epoch in range(nEpochs):
13        error = 0.0
14        for i in range(int(usingImages*0.8/batchSize)):
15            xs = X_train[i*batchSize:(i+1)*batchSize]
16            ys = y_train[i*batchSize:(i+1)*batchSize]
17            _, er = sess.run([optimizer,loss],feed_dict={x:xs,t:ys↵
                })
18            error += er
19            if (epoch+1)%printEvery == 0:
20                print('Loss in ',epoch+1,' epoch is ',error/(↵
                    usingImages*0.8))
21
22    prediction = tf.equal(tf.argmax(y),tf.argmax(t))
23    accuracy = tf.reduce_mean(tf.cast(prediction,"float"))
24    print("Accuracy validation:", sess.run(accuracy,{x: X_validate,↵
        t: y_validate}))
25    print("Accuracy Test:", sess.run(accuracy,{x: X_test, t: y_test↵
        })))
```

Code Output

Listing 1: Code

```
1 nEpochs = 20
2 keepRate = 0.8
3 start_time = time.time()
4 y, regLoss = convolutionalNeuralNetwork(x)
5 print("Y:::",y)
6 print(t)
7 trainNetwork()
8 print("--- %s seconds ---" % (time.time() - start_time))
```

Listing 2: Code output.

```
1 Tensor("MaxPool_0", shape=(?, 109, 89, 64), dtype=float32)
2 Tensor("MaxPool_1:0", shape=(?, 55, 45, 128), dtype=float32)
3 Tensor("MaxPool_2:0", shape=(?, 28, 23, 256), dtype=float32)
4 Tensor("MaxPool_3:0", shape=(?, 14, 12, 256), dtype=float32)
5 Tensor("Reshape:0", shape=(?, 43008), dtype=float32)
6 Tensor("Relu_4:0", shape=(?, 1024), dtype=float32)
7 Tensor("add_6:0", shape=(?, 2), dtype=float32)
8 Y::: Tensor("add_6:0", shape=(?, 2), dtype=float32)
9 Tensor("Placeholder_1:0", shape=(?, 2), dtype=float32)
10 Loss in 1 epoch is 3477590.3936
11 Loss in 2 epoch is 518266.9358
12 Loss in 3 epoch is 289935.3949
13 Loss in 4 epoch is 219638.7662
14 Loss in 5 epoch is 145500.0971
15 Loss in 6 epoch is 95223.83015
16 Loss in 7 epoch is 75723.7899227
17 Loss in 8 epoch is 66588.3607875
18 Loss in 9 epoch is 41361.5934953
19 Loss in 10 epoch is 42081.5419438
20 Loss in 11 epoch is 28284.8736906
21 Loss in 12 epoch is 20337.5216562
22 Loss in 13 epoch is 18180.2584531
23 Loss in 14 epoch is 17251.6861375
24 Loss in 15 epoch is 12881.9880367
```

```
25 Loss in 16 epoch is 12475.3094367
26 Loss in 17 epoch is 8886.34049531
27 Loss in 18 epoch is 8565.50805313
28 Loss in 19 epoch is 5204.7743375
29 Loss in 20 epoch is 4545.29912266
30 Accuracy training: 96.838420
31 Accuracy validation: 96.491309
32 Accuracy Test: 96.384021
33 --- 814.2028067111969 seconds ---
```

Results

As from the above output, accuracy and other parameters are as follows...

- Train accuracy is: 96.83
- Validation accuracy is: 96.49
- Test accuracy is: 96.38
- The learning rate is 0.001
- A drop-out rate of 0.2 is used, i.e., keepRate = 0.8.
- I used the default resolution as it is but decreasing the resolution has decreased the accuracy a little.
- I used 4 convolutional layers.
- $[Conv2 \rightarrow ReLU \rightarrow Pooling] \rightarrow [Conv2 \rightarrow ReLU \rightarrow Pooling] \rightarrow [Conv2 \rightarrow ReLU \rightarrow Pooling] \rightarrow [Conv2 \rightarrow ReLU \rightarrow Pooling] \rightarrow [FC \rightarrow FC] \rightarrow FC$
- In order to make the most out the loaded data, it is rotated by 15° .
- Since the output prediction is just 2 classes, I used Sigmoid instead of Softmax.