

# Linear Models for Regression

Sargur Srihari  
srihari@cedar.buffalo.edu

# Plan of Discussion

- Discuss supervised learning starting with regression
- Goal of regression:
  - predict value of one or more target variables  $t$
  - Given  $d$ -dimensional vector  $\mathbf{x}$  of input variables

# Topics in Linear Regression

- Regression Terminology
- Polynomial Curve Fitting with Scalar input
- Linear Model with  $D$  inputs
- Linear Basis Function Models
  - Gaussian Radial Basis Functions
- Maximum Likelihood and Least Squares
- Geometry of Least Squares
- Sequential Learning
- Regularized Least Squares

# Regression Terminology

- Regression
  - Predict a numerical value  $t$  given some input
    - Learning algorithm has to output function  $f: \mathbb{R}^n \rightarrow \mathbb{R}$ 
      - where  $n$  = no of input variables
  - Ex: expected claim amount an insured person will make (used to set insurance premiums) or prediction of future prices of securities
    - Also used for algorithmic trading
- Classification
  - If  $t$  value is a label (categories):  $f: \mathbb{R}^n \rightarrow \{1, \dots, k\}$
- Ordinal Regression
  - Discrete values, ordered categories

# Polynomial Curve Fitting with a Scalar

- Simplest form of regression:


- With a single input variable  $x$

- $$y(x, \mathbf{w}) = w_0 + w_1 x + w_2 x^2 + \dots + w_M x^M = \sum_{j=0}^M w_j x^j$$

$M$  is the order of the polynomial,

$x^j$  denotes  $x$  raised to the power  $j$ ,

Coefficients  $w_0, \dots, w_M$  are collectively denoted by vector  $\mathbf{w}$



Training data set  
N=10, Input  $x$ , target  $t$

- Task: Learn  $\mathbf{w}$  from training data  $D = \{(x_i, t_i)\}, i = 1, \dots, N$

- Can be done by minimizing an error function that minimizes the misfit between  $y(x, \mathbf{w})$  for any given  $\mathbf{w}$  and training data
  - One simple choice of error function is sum of squares of error between predictions  $y(x_n, \mathbf{w})$  for each data point  $x_n$  and corresponding target values  $t_n$  so that we minimize

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{y(x_n, \mathbf{w}) - t_n\}^2$$

- It is zero when function  $y(x, \mathbf{w})$  passes exactly through each training data point

# Regression with multiple inputs

- Generalization
  - Predict value of continuous target variable  $t$  given value of  $d$  input variables  $\mathbf{x}=[x_1, \dots, x_D]$
  - $t$  can also be a set of variables (multiple regression)
  - Linear functions of adjustable parameters
    - Specifically linear combinations of nonlinear functions of input variable
- Polynomial curve fitting is good only for:
  - Single input variable scalar  $x$
  - It cannot be easily generalized to several variables, as we will see

# Simplest Linear Model with $D$ inputs

- Regression with  $D$  input variables

$$y(\mathbf{x}, \mathbf{w}) = w_0 + w_1 x_1 + \dots + w_d x_d = \mathbf{w}^T \mathbf{x}$$

This differs from  
Linear Regression with one variable  
and Polynomial Reg with one variable

where  $\mathbf{x} = (x_1, \dots, x_D)^T$  are the input variables


- Called Linear Regression since it is a linear function of
  - parameters  $w_0, \dots, w_D$
  - input variables  $x_1, \dots, x_D$
- Significant limitation since it is a linear function of input variables
  - In the one-dimensional case this amounts a straight-line fit (degree-one polynomial)
  - $y(x, \mathbf{w}) = w_0 + w_1 x$

# Fitting a Regression Plane

- Assume  $t$  is a function of inputs  $x_1, x_2, \dots, x_D$   
Goal: find best linear regressor of  $t$  on all inputs

- Fitting a hyperplane through  $N$  input samples

- For  $D=2$ :



| $x_1$ | $x_2$ | $t$ |
|-------|-------|-----|
| 1     | 2     | 2   |
| 2     | 5     | 1   |
| 2     | 3     | 2   |
| 2     | 2     | 2   |
| 3     | 4     | 1   |
| 3     | 5     | 3   |
| 4     | 6     | 2   |
| 5     | 5     | 3   |
| 5     | 6     | 4   |
| 5     | 7     | 3   |
| 6     | 8     | 4   |
| 7     | 6     | 2   |
| 8     | 4     | 4   |
| 8     | 9     | 3   |
| 9     | 8     | 4   |

- Being a linear function of input variables imposes limitations on the model
  - Can extend class of models by considering fixed nonlinear functions of input variables



# Basis Functions

- In many applications, we apply some form of fixed-preprocessing, or feature extraction, to the original data variables
- If the original variables comprise the vector  $\mathbf{x}$ , then the features can be expressed in terms of basis functions  $\{\phi_j(\mathbf{x})\}$ 
  - By using nonlinear basis functions we allow the function  $y(\mathbf{x}, \mathbf{w})$  to be a nonlinear function of the input vector  $\mathbf{x}$ 
    - They are linear functions of parameters (gives them simple analytical properties), yet are nonlinear wrt input variables

# Linear Regression with $M$ Basis Functions

- Extended by considering nonlinear functions of input variables

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x})$$

- where  $\phi_j(\mathbf{x})$  are called Basis functions
- We now need  $M$  weights for basis functions instead of  $D$  weights for features
- With a dummy basis function  $\phi_0(\mathbf{x})=1$  corresponding to the bias parameter  $w_0$ , we can write

$$y(\mathbf{x}, \mathbf{w}) = \sum_{j=0}^{M-1} w_j \phi_j(\mathbf{x}) = \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x})$$

- where  $\mathbf{w} = (w_0, w_1, \dots, w_{M-1})$  and  $\boldsymbol{\phi} = (\phi_0, \phi_1, \dots, \phi_{M-1})^T$

- Basis functions allow non-linearity with  $D$  input variables

# Choice of Basis Functions

- Many possible choices for basis function:
  1. Polynomial regression
    - Good only if there is only one input variable
  2. Gaussian basis functions
  3. Sigmoidal basis functions
  4. Fourier basis functions
  5. Wavelets


# 1. Polynomial Basis for one variable

- Linear Basis Function Model

$$y(\mathbf{x}, \mathbf{w}) = \sum_{j=0}^{M-1} w_j \phi_j(\mathbf{x}) = \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x})$$

- Polynomial Basis (for single variable  $x$ )

$\phi_j(x) = x^j$  with degree  $M-1$  polynomial



- Disadvantage

- Global:

- changes in one region of input space affects others

- Difficult to formulate

- Number of polynomials increases exponentially with  $M$

- Can divide input space into regions

- use different polynomials in each region:
- equivalent to spline functions

# Can we use Polynomial with $D$ variables? (Not practical!)

- Consider (for a vector  $\mathbf{x}$ ) the basis:  $\phi_j(\mathbf{x}) = \|\mathbf{x}\|^j = \left[ \sqrt{x_1^2 + x_2^2 + \dots + x_d^2} \right]^j$ 
  - $\mathbf{x}=(2,1)$  and  $\mathbf{x}=(1,2)$  have the same squared sum, so it is unsatisfactory
  - Vector is being converted into a scalar value thereby losing information
- Better polynomial approach:
  - Polynomial of degree  $M-1$  has terms with variables taken none, one, two...  $M-1$  at a time.
  - Use multi-index  $j=(j_1, j_2, \dots, j_D)$  such that  $j_1 + j_2 + \dots + j_D \leq M-1$
  - For a quadratic ( $M=3$ ) with three variables ( $D=3$ )

$$y(\mathbf{x}, \mathbf{w}) = \sum_{(j_1, j_2, j_3)} w_j \phi_j(\mathbf{x}) = w_0 + w_{1,0,0}x_1 + w_{0,1,0}x_2 + w_{0,0,1}x_3 + w_{1,1,0}x_1x_2 + w_{1,0,1}x_1x_3 + w_{0,1,1}x_2x_3 + w_{2,0,0}x_1^2 + w_{0,2,0}x_2^2 + w_{0,0,2}x_3^2$$

- Number of quadratic terms is  $1 + D + D(D-1)/2 + D$
- For  $D=46$ , it is 1128
- Better to use Gaussian kernel, discussed next

## 2. Gaussian Radial Basis Functions

- Gaussian

$$\phi_j(x) = \exp\left(-\frac{(x - \mu_j)^2}{2\sigma^2}\right)$$

- Does not necessarily have a probabilistic interpretation
- Usual normalization term is unimportant
  - since basis function is multiplied by weight  $w_j$


- Choice of parameters

- $\mu_j$  govern the locations of the basis functions
  - Can be an arbitrary set of points within the range of the data
    - Can choose some representative data points
- $\sigma$  governs the spatial scale
  - Could be chosen from the data set e.g., average variance

- Several variables

- A Gaussian kernel would be chosen for each dimension
- For each  $j$  a different set of means would be needed– perhaps chosen from the data


$$\phi_j(\mathbf{x}) = \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_j)^t \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu}_j)\right)$$




# Biological Inspiration for RBFs

- Nervous system contains many examples

– Local receptive fields in visual cortex



- RBF Network




- Receptive fields overlap
- So there is usually more than one unit active
- But for given input, total no. of active units is small

# Tiling the input space

- Determining centers

- $k$ -means clustering

- Choose  $k$  cluster centers
    - Mark each training point as captured by cluster to which it is closest
    - Move each cluster center to mean of points it captured



- Determining variance  $\sigma^2$

Global:  $\sigma$  = mean distance between each unit  $j$  and its closest neighbor

P-nearest neighbor: set each  $\sigma_j$  so that there is certain overlap with P closest neighbors of unit  $j$




# 3. Sigmoidal Basis Function

- Sigmoid  $\phi_j(x) = \sigma\left(\frac{x - \mu_j}{s}\right)$  where  $\sigma(a) = \frac{1}{1 + \exp(-a)}$
- Equivalently,  $\tanh$  because it is related to logistic sigmoid by

$$\tanh(a) = 2\sigma(a) - 1$$

Logistic Sigmoid  
For different  $\mu_j$



## 4. Other Basis Functions

- Fourier
  - Expansion in sinusoidal functions
  - Infinite spatial extent
- Signal Processing
  - Functions localized in time and frequency
  - Called *wavelets*
    - Useful for lattices such as images and time series
- Further discussion independent of choice of basis including  $\phi(x) = x$

# Relationship between Maximum Likelihood and Least Squares


- Will show that Minimizing sum-of-squared errors is the same as maximum likelihood solution under a Gaussian noise model
- Target variable is a scalar  $t$  given by deterministic function  $y(x, \mathbf{w})$  with additive Gaussian noise  $\varepsilon$

$$t = y(x, \mathbf{w}) + \varepsilon$$

– which is a *zero-mean* Gaussian with *precision*  $\beta$

- Thus distribution of  $t$  is univariate normal:


$$p(t|x, \mathbf{w}, \beta) = N(t \mid \underbrace{y(x, \mathbf{w})}_{\text{mean}}, \underbrace{\beta^{-1}}_{\text{variance}})$$



# Likelihood Function

- Data set:
  - Input  $X = \{x_1, \dots, x_N\}$  with target  $t = \{t_1, \dots, t_N\}$
  - Target variables  $t_n$  are scalars forming a vector of size  $N$
- Likelihood of the target data
  - It is the probability of observing the data assuming they are independent
  - since  $p(t|x, w, \beta) = N(t | y(x, w), \beta^{-1})$
  - and  $y(x, w) = \sum_{j=0}^{M-1} w_j \phi_j(x) = w^T \phi(x)$

$$p(t | X, w, \beta) = \prod_{n=1}^N N(t_n | w^T \phi(x_n), \beta^{-1})$$



# Log-Likelihood Function

- Likelihood

$$p(\mathbf{t} | \mathbf{X}, \mathbf{w}, \beta) = \prod_{n=1}^N N(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1})$$

- Log-likelihood

$$\ln p(\mathbf{t} | \mathbf{w}, \beta) = \sum_{n=1}^N \ln N(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1})$$

- Using standard univariate Gaussian

$$N(x | \mu, \sigma^2) = \frac{1}{(2\pi\sigma^2)^{1/2}} \exp\left\{-\frac{1}{2\sigma^2}(x - \mu)^2\right\}$$

$$\ln p(\mathbf{t} | \mathbf{w}, \beta) = \frac{N}{2} \ln \beta - \frac{N}{2} \ln 2\pi - \beta E_D(\mathbf{w})$$

- Where

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left\{ t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \right\}^2$$

Sum-of-squares Error Function

- With Gaussian basis functions

$$\phi_j(\mathbf{x}) = \exp\left(\frac{(\mathbf{x} - \boldsymbol{\mu}_j)^T (\mathbf{x} - \boldsymbol{\mu}_j)}{2s^2}\right)$$

# Maximizing Log-Likelihood Function

- Log-likelihood

$$\begin{aligned}\ln p(t | \mathbf{w}, \beta) &= \sum_{n=1}^N \ln N(t_n | \mathbf{w}^T \phi(\mathbf{x}_n), \beta^{-1}) \\ &= \frac{N}{2} \ln \beta - \frac{N}{2} \ln 2\pi - \beta E_D(\mathbf{w})\end{aligned}$$

– where

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(x_n)\}^2$$

- Therefore, maximizing likelihood is equivalent to minimizing  $E_D(\mathbf{w})$

# Determining max likelihood solution

- The likelihood function has the form

$$\ln p(t | w, \beta) = \frac{N}{2} \ln \beta - \frac{N}{2} \ln 2\pi - \beta E_D(w)$$
$$E_D(w) = \frac{1}{2} \sum_{n=1}^N \left\{ t_n - w^T \phi(x_n) \right\}^2$$

- where

- We will show that the maximum likelihood solution has a closed form
  - Take derivative of  $\ln p(t | w, \beta)$  wrt  $w$  and set equal to zero and solve for  $w$ 
    - or equivalently just the derivative of  $E_D(w)$

# Gradient of Log-likelihood wrt $\mathbf{w}$

$$\nabla \ln p(\mathbf{t} | \mathbf{w}, \beta) = \beta \sum_{n=1}^N \left\{ t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n) \right\} \boldsymbol{\phi}(\mathbf{x}_n)^T$$

-which is obtained from log-likelihood expression and by using calculus result

$$\nabla_w \left[ -\frac{1}{2} (a - wb)^2 \right] = (a - wb)b$$

- Gradient is set to zero and we solve for  $\mathbf{w}$

$$0 = \sum_{n=1}^N t_n \boldsymbol{\phi}(\mathbf{x}_n) - \mathbf{w}^T \left( \sum_{n=1}^N \boldsymbol{\phi}(\mathbf{x}_n) \boldsymbol{\phi}(\mathbf{x}_n)^T \right)$$

as shown in next slide

- Second derivative will be negative making this a maximum



# Max Likelihood Solution for $\mathbf{w}$

- Solving for  $\mathbf{w}$  we obtain:

$$\mathbf{w}_{ML} = \Phi^+ \mathbf{t}$$

$X = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$  are samples  
(vectors of  $d$  variables)

$\mathbf{t} = \{t_1, \dots, t_N\}$  are targets (scalars)

where  $\Phi^+ = (\Phi^T \Phi)^{-1} \Phi^T$  is the Moore-Penrose pseudo inverse of the  $N \times M$  Design Matrix  $\Phi$  whose elements are given by  $\Phi_{nj} = \phi_j(\mathbf{x}_n)$

- Known as the normal equations for the least squares problem

**Design Matrix:**  
Rows correspond to  
 $N$  samples,  
Columns to  
 $M$  basis functions

$$\Phi = \begin{pmatrix} \phi_0(\mathbf{x}_1) & \phi_1(\mathbf{x}_1) & \dots & \phi_{M-1}(\mathbf{x}_1) \\ \phi_0(\mathbf{x}_2) & & & \\ & & & \\ \phi_0(\mathbf{x}_N) & & & \phi_{M-1}(\mathbf{x}_N) \end{pmatrix}$$

**Pseudo inverse:**

generalization of notion of matrix inverse  
to non-square matrices

If design matrix is square and invertible.  
then pseudo-inverse is same as inverse

$\phi_i(\mathbf{x}_n)$  are  $M$  basis functions, e.g., Gaussians centered on  $M$  data points

$$\phi_j(\mathbf{x}) = \exp\left[-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_j)^t \Sigma^{-1}(\mathbf{x} - \boldsymbol{\mu}_j)\right]$$

# Design Matrix

$\leftarrow M$  Basis functions  $\rightarrow$

$$\Phi = \begin{pmatrix} \phi_0(x_1) & \phi_1(x_1) & \dots & \phi_{M-1}(x_1) \\ \phi_0(x_2) & & & \\ \vdots & & & \\ \phi_0(x_N) & & & \phi_{M-1}(x_N) \end{pmatrix}$$

$\uparrow N$  Data  $\downarrow$

Represented as  $N$ -dim vectors

$\phi_0 \quad \phi_1 \quad \phi_{M-1}$

The diagram illustrates the design matrix  $\Phi$  as a collection of  $N$ -dimensional vectors. The matrix is defined by basis functions  $\phi_0, \phi_1, \dots, \phi_{M-1}$  applied to data points  $x_1, x_2, \dots, x_N$ . The matrix is shown as a grid of elements. A horizontal arrow above the matrix indicates  $M$  basis functions. A vertical arrow to the right of the matrix indicates  $N$  data points. Below the matrix, the text 'Represented as  $N$ -dim vectors' is shown. Three arrows point from a common point below to the first, second, and last columns of the matrix, labeled  $\phi_0$ ,  $\phi_1$ , and  $\phi_{M-1}$  respectively.

# What is the role of Bias parameter $w_0$ ?

- Sum-of squares error function is:

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left\{ t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n) \right\}^2$$

- Substituting:  $y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x})$  we get:

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left\{ t_n - w_0 - \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x}_n) \right\}^2$$

- Setting derivatives wrt  $w_0$  equal to zero and solving for  $w_0$

$$w_0 = \bar{t} - \sum_{j=1}^{M-1} w_j \bar{\phi}_j$$

where  $\bar{t} = \frac{1}{N} \sum_{n=1}^N t_n$  and  $\bar{\phi}_j = \frac{1}{N} \sum_{n=1}^N \phi_j(\mathbf{x}_n)$

- First term is average of the  $N$  values of  $t$
- Second term is weighted sum of the average basis function values over  $N$  samples
- Thus bias  $w_0$  compensates for difference between average target values and weighted sum of averages of basis function values

# Maximum Likelihood for precision $\beta$

- We have determined m.l.e. solution for  $\mathbf{w}$  using a probabilistic formulation

- $p(t|\mathbf{x}, \mathbf{w}, \beta) = N(t|y(\mathbf{x}, \mathbf{w}), \beta^{-1})$

- With log-likelihood


$$\ln p(\mathbf{t}|\mathbf{w}, \beta) = \frac{N}{2} \ln \beta - \frac{N}{2} \ln 2\pi - \beta E_D(\mathbf{w})$$

$$\mathbf{w}_{ML} = \Phi^+ \mathbf{t}$$

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{ t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \}^2$$

- Taking gradient wrt  $\beta$  gives


$$\frac{1}{\beta_{ML}} = \frac{1}{N} \sum_{n=1}^N \{ t_n - \mathbf{w}_{ML}^T \phi(\mathbf{x}_n) \}^2$$



- Thus Inverse of the noise precision gives  
*Residual variance of the target values  
around the regression function*

# Geometry of Least Squares

- Geometrical Interpretation of Least Squares Solution instructive
- Consider  $N$ -dim space with axes  $t_n$   
so that  $\mathbf{t} = (t_1, \dots, t_N)^T$  is a vector in this space
- Each basis  $\phi_j(\mathbf{x}_n)$  evaluated at  $N$  points can also be represented as a vector in the same space
- $\phi_j$  corresponds to  $j^{th}$  column of  $\Phi$ , whereas  $\phi(\mathbf{x}_n)$  corresponds to the  $n^{th}$  row of  $\Phi$
- If the no of basis functions is smaller than the no of data points  
– i.e.,  $M < N$  then the  $M$  vectors  $\phi_j(\mathbf{x}_n)$  will span linear subspace  $S$  of dim  $M$
- Define  $\mathbf{y}$  to be an  $N$ -dim vector whose  $n^{th}$  element is  $y(\mathbf{x}_n, \mathbf{w})$
- Sum-of-squares error is equal to squared Euclidean distance between  $\mathbf{y}$  and  $\mathbf{t}$
- Solution  $\mathbf{w}$  corresponds to  $\mathbf{y}$  that lies in subspace  $S$  that is closest to  $\mathbf{t}$ 
  - Corresponds to orthogonal projection of  $\mathbf{t}$  onto  $S$



# Difficulty of Direct solution

- Direct solution of normal equations

$$\boxed{w_{ML} = \Phi^+ t} \quad \boxed{\Phi^+ = (\Phi^T \Phi)^{-1} \Phi^T}$$


- This direct solution can lead to numerical difficulties
  - When  $\Phi^T \Phi$  is close to singular (determinant=0)
  - When two basis functions are collinear parameters can have large magnitudes
- Not uncommon with real data sets
- Can be addressed using
  - Singular Value Decomposition
  - Addition of regularization term ensures matrix is non-singular

# Method of Gradient Descent


- Criterion  $f(x)$  is minimized by moving from the current solution in direction of the negative of gradient
- Steepest descent proposes a new point

$$x' = x - \varepsilon \nabla_x f(x)$$

- where  $\varepsilon$  is the learning
- rate, a positive scalar.
- Set to a small constant.



# Direction of Steepest Descent





# Gradient Descent for Regression

- Error function  $E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{ t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \}^2$  sums over data
  - Denoting  $E_D(\mathbf{w}) = \sum_n E_n$ , update parameter vector  $\mathbf{w}$  using

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E_n$$

- where  $\tau$  is the iteration number and  $\eta$  is a learning rate parameter

- Substituting for the derivative

$$\nabla E_n = - \sum_{n=1}^N \{ t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \} \phi(\mathbf{x}_n)^T$$

where  $\phi_n = \phi(\mathbf{x}_n)$

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} + \eta (t_n - \mathbf{w}^{(\tau)T} \phi_n) \phi_n$$

- $\mathbf{w}$  is initialized to some starting vector  $\mathbf{w}^{(0)}$
- $\eta$  chosen with care to ensure convergence
- Known as *Least Mean Squares* Algorithm

# Sequential (On-line) Learning

- Maximum likelihood solution is

$$\mathbf{w}_{ML} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

- It is a batch technique
  - Processing entire training set in one go
- It is computationally expensive for large data sets
  - Due to huge  $N \times M$  Design matrix  $\Phi$
- Solution is to use a sequential algorithm where samples are presented one at a time (or a minibatch at a time)
- Called stochastic gradient descent

# Regularized Least Squares

- As model complexity increases, e.g., degree of polynomial or no. of basis functions, then it is likely that we overfit
- One way to control overfitting is not to limit complexity but to add a regularization term to the error function
- Error function to minimize takes the form

$$E(\mathbf{w}) = E_D(\mathbf{w}) + \lambda E_W(\mathbf{w})$$

- where  $\lambda$  is the *regularization coefficient* that controls relative importance of data-dependent error  $E_D(\mathbf{w})$  and regularization term  $E_W(\mathbf{w})$

# Simplest Regularizer is weight decay

- Regularized least squares is

$$E(\mathbf{w}) = E_D(\mathbf{w}) + \lambda E_W(\mathbf{w})$$

- Simple form of regularization term is

$$E_W(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w}$$

- Thus total error function becomes

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left\{ t_n - \mathbf{w}^T \phi(x_n) \right\}^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

- This regularizer is called *weight decay*
  - because in sequential learning weight values decay towards zero unless supported by data
- Also, the error function remains a *quadratic function* of  $\mathbf{w}$ , so exact minimizer found in closed form

# Closed-form Solution with Regularizer

- Error function with *quadratic regularizer* is,

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left\{ t_n - \mathbf{w}^T \phi(x_n) \right\}^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$


- Its exact minimizer can be found in closed form
  - By setting gradient wrt  $\mathbf{w}$  to zero and solving for  $\mathbf{w}$

$$\mathbf{w} = (\lambda \mathbf{I} + \Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

- This is a simple extension of the least squared solution

$$\mathbf{w}_{ML} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

# Geometric Interpretation of Regularizer



Any value of  $w$  on contour gives same error


In *unregularized* case:  
we are trying to find  $w$  that minimizes

$$E_D(w) = \frac{1}{2} \sum_{n=1}^N \left\{ t_n - w^T \phi(x_n) \right\}^2$$

In *regularized* case:  
choose that value of  $w$  subject to the constraint

$$\sum_{j=1}^M |w_j|^2 \leq \eta$$


We don't want the weights to become too large  
The two approaches related by Lagrange multipliers



$$E(w) = \frac{1}{2} \sum_{n=1}^N \left\{ t_n - w^T \phi(x_n) \right\}^2 + \frac{\lambda}{2} w^T w$$

# Minimization of Unregularized Error subject to constraint

- Blue: Contours of unregularized error function
- Constraint region
- $w^*$  is optimum value



# A more general regularizer

- Regularized Error

$$\frac{1}{2} \sum_{n=1}^N \left\{ t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \right\}^2 + \frac{\lambda}{2} \sum_{j=1}^M |w_j|^q$$

- Where  $q=2$  corresponds to the *quadratic* regularizer  
 $q=1$  is known as *lasso*
- Lasso has the property that if  $\lambda$  is sufficiently large some of the coefficients  $w_j$  are driven to zero leading to a sparse model in which the corresponding basis functions play no role



# Contours of regularization term


$$\frac{1}{2} \sum_{n=1}^N \left\{ t_n - \mathbf{w}^T \phi(\mathbf{x}_n) \right\}^2 + \frac{\lambda}{2} \sum_{j=1}^M |w_j|^q$$

- Contours of regularization term  $|w_j|^q$  for values of  $q$


Space of  $w_1, w_2$

Any choice along the contour has the same value of  $\mathbf{w}$


$$\sqrt{w_1} + \sqrt{w_2} = \text{const}$$




$$w_1 + w_2 = \text{const}$$



$$w_1^2 + w_2^2 = \text{const}$$




$$w_1^4 + w_2^4 = \text{const}$$




# Sparsity with Lasso constraint

- With  $q=1$  and  $\lambda$  is sufficiently large, some of the coefficients  $w_j$  are driven to zero
- Leads to a sparse model
  - where corresponding basis functions play no role
- Origin of sparsity is illustrated here:

Quadratic solution where  $w_1^*$  and  $w_0^*$  are nonzero



Minimization with Lasso Regularizer  
A sparse solution with  $w_1^*=0$



# Regularization: Conclusion

- Regularization allows
  - complex models to be trained on small data sets
  - without severe over-fitting
- It limits model complexity
  - i.e., how many basis functions to use?
- Problem of limiting complexity is shifted to
  - one of determining suitable value of regularization coefficient

# Linear Regression Summary

- Linear Regression with  $M$  basis functions:

$$y(\mathbf{x}, \mathbf{w}) = \sum_{j=0}^{M-1} w_j \phi_j(\mathbf{x}) = \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x})$$

$$\phi_j(\mathbf{x}) = \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_j)^T \boldsymbol{\Sigma}^{-1}(\mathbf{x} - \boldsymbol{\mu}_j)\right)$$

- Objective Function *without/with* regularization is

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left\{ t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n) \right\}^2$$

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \left\{ t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n) \right\}^2 + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w}$$

- Closed-form ML solution is:

$$\mathbf{w}_{ML} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

$$\mathbf{w}_{ML} = (\lambda I + \Phi^T \Phi)^{-1} \Phi^T \mathbf{t}$$

$$\Phi = \begin{pmatrix} \phi_0(\mathbf{x}_1) & \phi_1(\mathbf{x}_1) & \dots & \phi_{M-1}(\mathbf{x}_1) \\ \phi_0(\mathbf{x}_2) & & & \\ & & & \\ \phi_0(\mathbf{x}_N) & & & \phi_{M-1}(\mathbf{x}_N) \end{pmatrix}$$

- Gradient Descent:  $\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E_n$

$$\nabla E_n = - \sum_{n=1}^N \left\{ t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n) \right\} \boldsymbol{\phi}(\mathbf{x}_n)^T$$

$$\nabla E_n = \left[ - \sum_{n=1}^N \left\{ t_n - \mathbf{w}^T \boldsymbol{\phi}(\mathbf{x}_n) \right\} \boldsymbol{\phi}(\mathbf{x}_n)^T \right] + \lambda \mathbf{w}$$

# Returning to LeToR Problem

- Try:
- Several Basis Functions
- Quadratic Regularization
- Express results as  $E_{RMS}$ 
  - rather than as squared error  $E(\mathbf{w}^*)$  or as Error Rate with thresholded results

$$E_{RMS} = \sqrt{2E(\mathbf{w}^*)/N}$$

# Multiple Outputs

- Several target variables  $\mathbf{t} = (t_1, \dots, t_K)$   $K > 1$
- Can be treated as multiple ( $K$ ) independent regression problems
  - Different basis functions for each component of  $\mathbf{t}$
- More common solution: same set of basis functions to model all components of target vector  $\mathbf{y}(\mathbf{x}, \mathbf{w}) = \mathbf{W}^T \boldsymbol{\phi}(\mathbf{x})$ 
  - where  $\mathbf{y}$  is a  $K$ -dim column vector,  $\mathbf{W}$  is a  $M \times K$  matrix of weights and  $\boldsymbol{\phi}(\mathbf{x})$  is a  $M$ -dimensional column vector with elements  $\phi_j(\mathbf{x})$

# Solution for Multiple Outputs

- Set of observations  $\mathbf{t}_1, \dots, \mathbf{t}_N$  are combined into a matrix  $\mathbf{T}$  of size  $N \times K$  such that the  $n^{\text{th}}$  row is given by  $\mathbf{t}_n^T$
- Combine input vectors  $\mathbf{x}_1, \dots, \mathbf{x}_N$  into matrix  $\mathbf{X}$
- Log-likelihood function is maximized
- Solution is similar:  $\mathbf{W}_{\text{ML}} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{T}$