

Numerical Computation

Sargur N. Srihari

srihari@cedar.buffalo.edu

This is part of lecture slides on [Deep Learning](http://www.cedar.buffalo.edu/~srihari/CSE676):
<http://www.cedar.buffalo.edu/~srihari/CSE676>

Topics

- Overflow and Underflow
- Poor Conditioning
- Gradient-based Optimization
- Stationary points, Local minima
- Second Derivative
- Convex Optimization
- Lagrangian

Acknowledgements: Goodfellow, Bengio, Courville, Deep Learning, MIT Press, 2016

Overview

- ML algorithms usually require a high amount of numerical computation
 - To update estimate of solutions iteratively
 - not analytically derive formula providing expression
- Common operations:
 - Optimization
 - Determine maximum or minimum of a function
 - Solving system of linear equations
- Just evaluating a mathematical function of real numbers with finite memory can be difficult

Overflow and Underflow

- Problem caused by representing real numbers with finite bit patterns
 - For almost all real numbers we encounter approximations
- Although a rounding error it compounds across many operations and algorithm will fail
 - Numerical errors
 - Underflow: when nos close to zero are rounded to zero
 - $\log 0$ is $-\infty$ (which becomes not-number for further operations)
 - Overflow: when nos with large magnitude are approximated as $-\infty$ or $+\infty$ (Again become not-no.)

Function needing stabilization for Over/Underflow

- Softmax probabilities in multinoulli

$$\text{softmax}(\mathbf{x})_i = \frac{\exp(x_i)}{\sum_{j=1}^n \exp(x_j)}$$

- Consider when all x_i are equal to some c . Then all probabilities must equal $1/n$. This may not happen

- When c is a large negative; denominator $\rightarrow 0$, result undefined underflow

- When c is large positive, $\exp(c)$ will overflow

- Circumvented using $\text{softmax}(\mathbf{z})$ where $\mathbf{z} = \mathbf{x} - \max_i x_i$

- Another problem: underflow in numerator can cause $\log \text{softmax}(\mathbf{x})$ to be $-\infty$

- Same trick can be used as for softmax

Dealing with numerical considerations

- Developers of low-level libraries should take this into consideration
- ML libraries should be able to provide such stabilization
 - Theano for Deep Learning detects and provides this

Poor Conditioning

- Conditioning refers to how rapidly a function changes with a small change in input
- Rounding errors can rapidly change the output
- Consider $f(\mathbf{x}) = A^{-1}\mathbf{x}$
 - $A \in \mathbb{R}^{n \times n}$ has a eigendecomposition
 - Its condition no. is $\max_{i,j} \left| \frac{\lambda_i}{\lambda_j} \right|$, i.e. ratio of largest to smallest eigenvalue
 - When this large, the output is very sensitive to input error
 - Poorly conditioned matrices amplify pre-existing errors when we multiply by its inverse

Gradient-Based Optimization

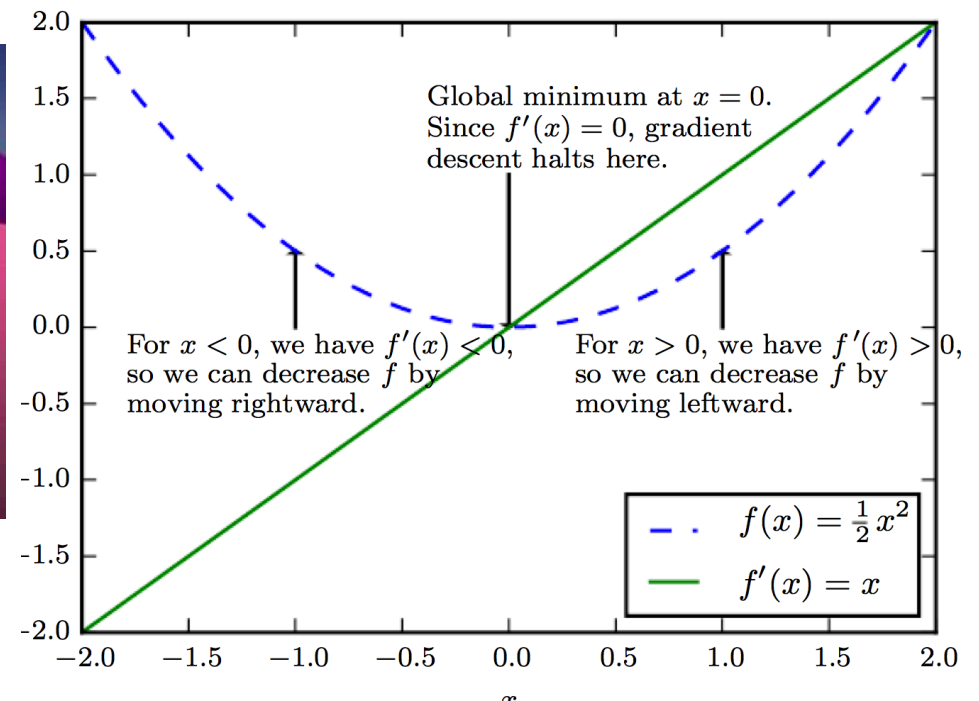
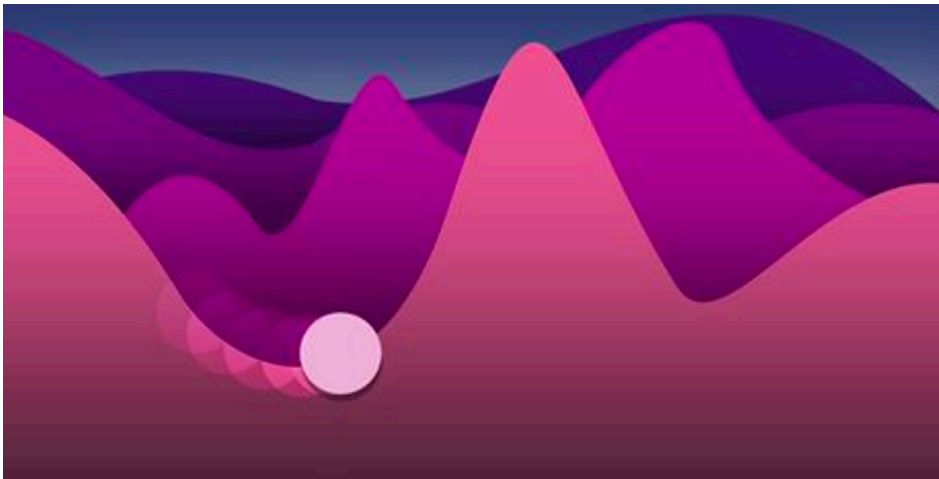
- Most ML algorithms involve optimization
- Minimize/maximize a function $f(\mathbf{x})$ by altering \mathbf{x}
 - Usually stated a minimization
 - Maximization accomplished by minimizing $-f(\mathbf{x})$
- $f(\mathbf{x})$ referred to as objective function or criterion
 - In minimization also referred to as loss function cost, or error
 - Example is linear least squares $f(x) = \frac{1}{2} ||Ax - b||^2$
 - Denote optimum value by $\mathbf{x}^* = \operatorname{argmin} f(\mathbf{x})$

Calculus in Optimization

- Suppose function $y=f(x)$, x , y real nos.
 - Derivative of function denoted: $f'(x)$ or as dy/dx
 - Derivative $f'(x)$ gives the slope of $f(x)$ at point x
 - It specifies how to scale a small change in input to obtain a corresponding change in the output:
$$f(x + \varepsilon) \approx f(x) + \varepsilon f'(x)$$
 - It tells how you make a small change in input to make a small improvement in y
 - We know that $f(x - \varepsilon \text{ sign}(f'(x)))$ is less than $f(x)$ for small ε . Thus we can reduce $f(x)$ by moving x in small steps with opposite sign of derivative
 - This technique is called *gradient descent* (Cauchy 1847)


Gradient Descent Illustrated

- For $x > 0$, $f(x)$ increases with x and $f'(x) > 0$
- For $x < 0$, $f(x)$ decreases with x and $f'(x) < 0$
- Use $f'(x)$ to follow function downhill
- Reduce $f(x)$ by going in direction opposite sign of derivative $f'(x)$



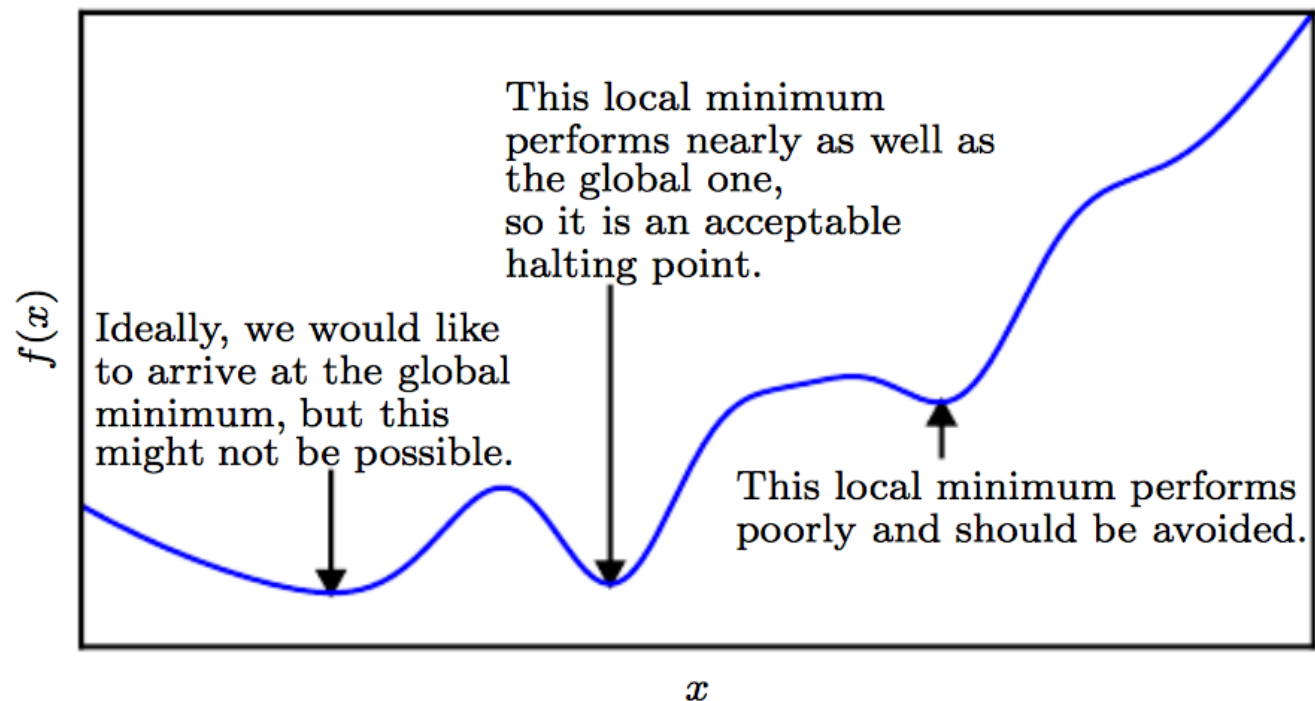
Stationary points, Local Optima

- When $f'(x)=0$ derivative provides no information about direction of move
- Points where $f'(x)=0$ are known as *stationary* or critical points
 - Local minimum/maximum: a point where $f(x)$ lower/higher than all its neighbors
 - Saddle Points: neither maxima nor minima



Presence of multiple minima

- Optimization algorithms may fail to find global minimum
- Generally accept such solutions



Minimizing with multiple inputs

- We often minimize functions with multiple inputs: $f: R^n \rightarrow R$
- For minimization to make sense there must still be only one (scalar) output

Functions with multiple inputs

- Need partial derivatives
- $\frac{\partial}{\partial x_i} f(\mathbf{x})$ measures how f changes as only variable x_i increases at point \mathbf{x}
- Gradient generalizes notion of derivative where derivative is wrt a vector
- Gradient is vector containing all of the partial derivatives denoted $\nabla_{\mathbf{x}} f(\mathbf{x})$
 - Element i of the gradient is the partial derivative of f wrt x_i
 - Critical points are where every element of the gradient is equal to zero

Directional Derivative

- Directional derivative in direction \mathbf{u} (a unit vector) is the slope of function f in direction \mathbf{u}
 - This evaluates to $\mathbf{u}^T \nabla_x f(\mathbf{x})$
- To minimize f find direction in which f decreases the fastest
 - Do this using $\min_{\mathbf{u}, \mathbf{u}^T \mathbf{u}=1} \mathbf{u}^T \nabla_x f(\mathbf{x}) = \min_{\mathbf{u}, \mathbf{u}^T \mathbf{u}=1} \|\mathbf{u}\|_2 \|\nabla_x f(\mathbf{x})\|_2 \cos \theta$
 - where θ is angle between \mathbf{u} and the gradient
 - Substitute $\|\mathbf{u}\|_2=1$ and ignore factors that not depend on \mathbf{u} this simplifies to $\min_{\mathbf{u}} \cos \theta$
 - This is minimized when \mathbf{u} points in direction opposite to gradient
 - In other words, the *gradient points directly uphill, and the negative gradient points directly downhill*

Method of Gradient Descent

- The gradient points directly uphill, and the negative gradient points directly downhill
- Thus we can decrease f by moving in the direction of the negative gradient
 - This is known as the method of steepest descent or gradient descent
- Steepest descent proposes a new point

$$\mathbf{x}' = \mathbf{x} - \varepsilon \nabla_{\mathbf{x}} f(\mathbf{x})$$

- where ε is the learning rate, a positive scalar. Set to a small constant.

Choosing ε : Line Search

- We can choose ε in several different ways
- Popular approach: set ε to a small constant
- Another approach is called *line search*:
- Evaluate $f(\mathbf{x} - \varepsilon \nabla_{\mathbf{x}} f(\mathbf{x}))$ for several values of ε and choose the one that results in smallest objective function value

Ex: Gradient Descent on Least Squares

- Criterion to minimize

$$f(\mathbf{x}) = \frac{1}{2} || A\mathbf{x} - \mathbf{b} ||^2$$

- Least squares regression

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N \{t_n - \mathbf{w}^T \phi(x_n)\}^2$$

- The gradient is

$$\nabla_{\mathbf{x}} f(\mathbf{x}) = A^T (A\mathbf{x} - \mathbf{b}) = A^T A\mathbf{x} - A^T \mathbf{b}$$

- Gradient Descent algorithm is

1. Set step size ε , tolerance δ to small, positive nos.

2. *while* $|| A^T A\mathbf{x} - A^T \mathbf{b} ||_2 > \delta$ *do*

$$\mathbf{x} \leftarrow \mathbf{x} - \varepsilon (A^T A\mathbf{x} - A^T \mathbf{b})$$

3. *end while*

Convergence of Steepest Descent

- Steepest descent converges when every element of the gradient is zero
 - In practice, very close to zero
- We may be able to avoid iterative algorithm and jump to the critical point by solving the equation $\nabla_x f(\mathbf{x}) = 0$ for \mathbf{x}

Generalization to discrete spaces

- Gradient descent is limited to continuous spaces
- Concept of repeatedly making the best small move can be generalized to discrete spaces
- Ascending an objective function of discrete parameters is called *hill climbing*

Beyond Gradient: Jacobian and Hessian matrices

- Sometimes we need to find all derivatives of a function whose input and output are both vectors
- If we have function $f: R^m \rightarrow R^n$
 - Then the matrix of partial derivatives is known as the Jacobian matrix J defined as

$$J_{i,j} = \frac{\partial}{\partial x_j} f(x)_i$$

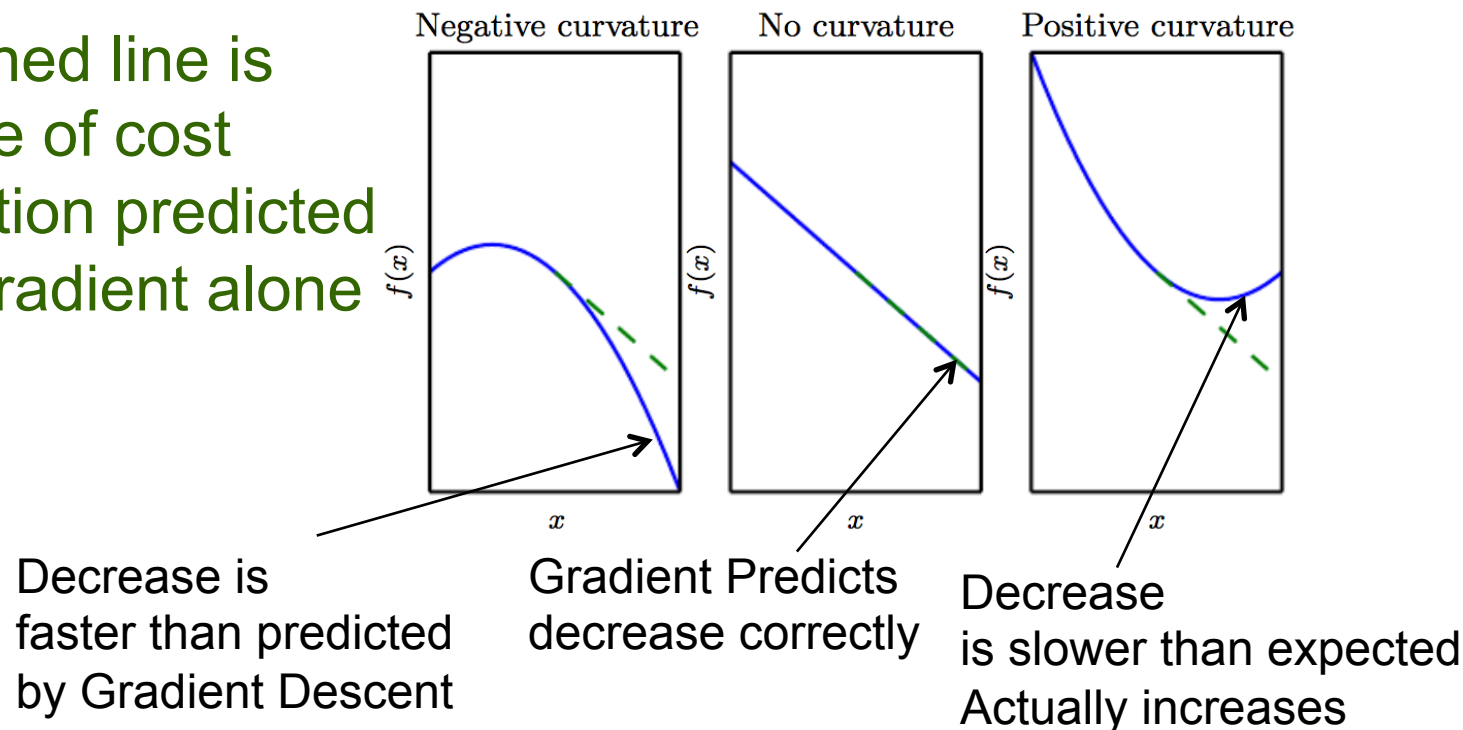
Second derivative

- Derivative of a derivative
- For a function $f: R^n \rightarrow R$ the derivative wrt x_i of the derivative of f wrt x_j is denoted as $\frac{\partial^2}{\partial x_i \partial x_j} f$
- In a single dimension we can denote $\frac{\partial^2}{\partial x^2} f$ by $f''(x)$
- Tells us how the first derivative will change as we vary the input
- This important as it tells us whether a gradient step will cause as much of an improvement as based on gradient alone

Second derivative measures curvature

- Derivative of a derivative
- Quadratic functions with different curvatures

Dashed line is
value of cost
function predicted
by gradient alone



Hessian

- Second derivative with many dimensions
- $H(f)(x)$ is defined as
$$H(f)(x)_{i,j} = \frac{\partial^2}{\partial x_i \partial x_j} f(x)$$
- Hessian is the Jacobian of the gradient
- Hessian matrix is symmetric, i.e., $H_{i,j} = H_{j,i}$
 - anywhere that the second partial derivatives are continuous
- So the Hessian matrix can be decomposed into a set of real eigenvalues and an orthogonal basis of eigenvectors
 - Eigenvalues of H are useful to determine learning rate as seen in next two slides

Role of eigenvalues of Hessian

- Second derivative in direction d is $d^T H d$
 - If d is an eigenvector, second derivative in that direction is given by its eigenvalue
 - For other directions, weighted average of eigenvalues (weights of 0 to 1, with eigenvectors with smallest angle with d receiving more value)
- Maximum eigenvalue determines maximum second derivative and minimum eigenvalue determines minimum second derivative

Learning rate from Hessian

- Taylor's series of $f(\mathbf{x})$ around current point $\mathbf{x}^{(0)}$

$$f(\mathbf{x}) \approx f(\mathbf{x}^{(0)}) + (\mathbf{x} - \mathbf{x}^{(0)})^T \mathbf{g} + \frac{1}{2} (\mathbf{x} - \mathbf{x}^{(0)})^T H (\mathbf{x} - \mathbf{x}^{(0)})$$

- where \mathbf{g} is the gradient and H is the Hessian at $\mathbf{x}^{(0)}$
- If we use learning rate ε the new point \mathbf{x} is given by $\mathbf{x}^{(0)} - \varepsilon \mathbf{g}$. Thus we get

$$f(\mathbf{x}^{(0)} - \varepsilon \mathbf{g}) \approx f(\mathbf{x}^{(0)}) - \varepsilon \mathbf{g}^T \mathbf{g} + \frac{1}{2} \varepsilon^2 \mathbf{g}^T H \mathbf{g}$$

- There are three terms:
 - original value of f ,
 - expected improvement due to slope, and
 - correction to be applied due to curvature
- Solving for step size when correction is least gives

$$\varepsilon^* \approx \frac{\mathbf{g}^T \mathbf{g}}{\mathbf{g}^T H \mathbf{g}}$$

Second Derivative Test: Critical Points


- On a critical point $f'(x)=0$
- When $f''(x)>0$ the first derivative $f'(x)$ increases as we move to the right and decreases as we move left
- We conclude that x is a local minimum
- For local maximum, $f'(x)=0$ and $f''(x)<0$
- When $f''(x)=0$ test is inconclusive: x may be a saddle point or part of a flat region

Multidimensional Second derivative test

- In multiple dimensions, we need to examine second derivatives of all dimensions
- Eigendecomposition generalizes the test
- Test eigenvalues of Hessian to determine whether critical point is a local maximum, local minimum or saddle point
- When H is positive definite (all eigenvalues are positive) the point is a local minimum
- Similarly negative definite implies a maximum

Saddle point

- Contains both positive and negative curvature
- Function is $f(\mathbf{x}) = x_1^2 - x_2^2$



- Along axis x_1 , function curves upwards: this axis is an eigenvector of H and has a positive value
- Along x_2 , function curves downwards; its direction is an eigenvector of H with negative eigenvalue
- At a saddle point eigen values are both positive and negative

Inconclusive Second Derivative Test

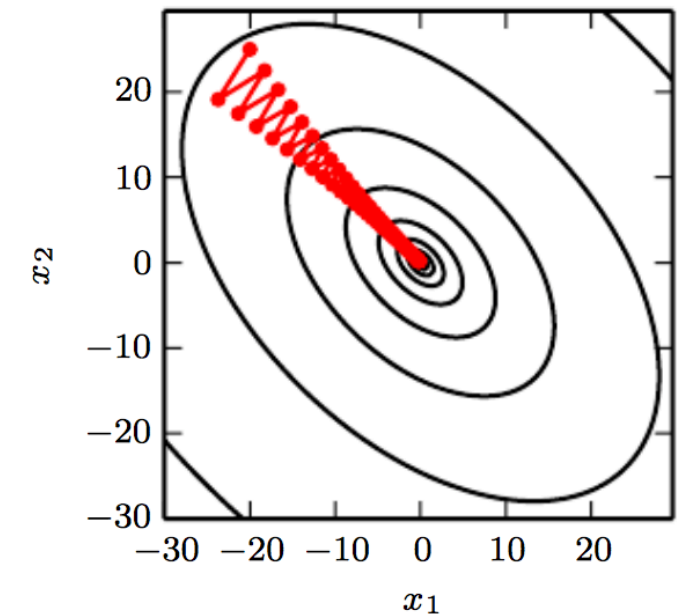
- Multidimensional second derivative test can be inconclusive just like univariate case
- Test is inconclusive when all non-zero eigenvalues have same sign but at least one value is zero
 - since univariate second derivative test is inconclusive in cross-section corresponding to zero eigenvalue

Poor Condition Number

- There are different second derivatives in each direction at a single point
- Condition number of H e.g., $\lambda_{max}/\lambda_{min}$ measures how much they differ
 - Gradient descent performs poorly when H has a poor condition no.
 - Because in one direction derivative increases rapidly while in another direction it increases slowly
 - Step size must be small so as to avoid overshooting the minimum, but it will be too small to make progress in other directions with less curvature

Gradient Descent without H

- H with condition no, 5
 - Direction of most curvature has five times more curvature than direction of least curvature
- Due to small step size
Gradient descent wastes time
- Algorithm based on Hessian
can predict that steepest
descent is not promising



Newton's method uses Hessian

- Another second derivative method
 - Using Taylor's series of $f(\mathbf{x})$ around current $\mathbf{x}^{(0)}$

$$f(\mathbf{x}) \approx f(\mathbf{x}^{(0)}) + (\mathbf{x} - \mathbf{x}^{(0)})^T \nabla_{\mathbf{x}} f(\mathbf{x}^{(0)}) + \frac{1}{2} (\mathbf{x} - \mathbf{x}^{(0)})^T H(f)(\mathbf{x} - \mathbf{x}^{(0)}) (\mathbf{x} - \mathbf{x}^{(0)})$$

- solve for the critical point of this function to give

$$\mathbf{x}^* = \mathbf{x}^{(0)} - H(f)(\mathbf{x}^{(0)})^{-1} \nabla_{\mathbf{x}} f(\mathbf{x}^{(0)})$$

- When f is a quadratic (positive definite) function use solution to jump to the minimum function directly
 - When not quadratic apply solution iteratively
- Can reach critical point much faster than gradient descent
 - But useful only when nearby point is a minimum

Summary of Gradient Methods

- First order optimization algorithms: those that use only the gradient
- Second order optimization algorithms: use the Hessian matrix such as Newton's method
- Family of functions used in ML is complicated, so optimization is more complex than in other fields
 - No guarantees
- Some guarantees by using Lipschitz continuous functions, $\boxed{\|f(\mathbf{x}) - f(\mathbf{y})\| \leq L \|\mathbf{x} - \mathbf{y}\|_2}$
 - with Lipschitz constant L

Convex Optimization

- Applicable only to convex functions—functions which are well-behaved,
 - e.g., lack saddle points and all local minima are global minima
- For such functions, Hessian is positive semi-definite everywhere
- Many ML optimization problems, particularly deep learning, cannot be expressed as convex optimization

Constrained Optimization

- We may wish to optimize $f(\mathbf{x})$ when the solution \mathbf{x} is constrained to lie in set S
 - Such values of \mathbf{x} are feasible solutions
- Often we want a solution that is small, such as $\|\mathbf{x}\| \leq 1$
- Simple approach: modify gradient descent taking constraint into account (using Lagrangian formulation)

Ex: Least squares with Lagrangian

- We wish to minimize
$$f(\mathbf{x}) = \frac{1}{2} || A\mathbf{x} - \mathbf{b} ||^2$$
- Subject to constraint $\mathbf{x}^T \mathbf{x} \leq 1$
- We introduce the Lagrangian
$$L(\mathbf{x}, \lambda) = f(\mathbf{x}) + \lambda(\mathbf{x}^T \mathbf{x} - 1)$$
 - And solve the problem
$$\min_{\mathbf{x}} \max_{\lambda, \lambda \geq 0} L(\mathbf{x}, \lambda)$$
- For the unconstrained problem (no Lagrangian) the smallest norm solution is $\mathbf{x} = A^+ \mathbf{b}$
 - If this solution is not feasible, differentiate Lagrangian wrt \mathbf{x} to obtain $A^T A \mathbf{x} - A^T \mathbf{b} + 2\lambda \mathbf{x} = 0$
 - Solution takes the form $\mathbf{x} = (A^T A + 2\lambda I)^{-1} A^T \mathbf{b}$
 - Choosing λ : continue solving linear equation and increasing λ until \mathbf{x} has the correct norm

Generalized Lagrangian: KKT

- More sophisticated than Lagrangian
- Karush-Kuhn-Tucker is a very general solution to constrained optimization
- While Lagrangian allows equality constraints, KKT allows both equality and inequality constraints
- To define a generalized Lagrangian we need to describe S in terms of equalities and inequalities

Generalized Lagrangian

- Set S is described in terms of m functions $g(i)$ and n functions $h(j)$ so that

$$S = \left\{ \mathbf{x} \mid \forall i, g^{(i)}(\mathbf{x}) = 0 \text{ and } \forall j, h^{(j)}(\mathbf{x}) \leq 0 \right\}$$

– Functions of g are equality constraints and functions of h are inequality constraints

- Introduce new variables λ_i and α_j for each constraint (called KKT multipliers) giving the generalized Lagrangian

$$L(\mathbf{x}, \lambda, \alpha) = f(\mathbf{x}) + \sum_i \lambda_i g^{(i)}(\mathbf{x}) + \sum_j \alpha_j h^{(j)}(\mathbf{x})$$

- We can now solve the unconstrained optimization problem

Gradient

- Essential role of calculus

