

Sequence Modeling: Recurrent and Recursive Nets

Sargur Srihari
srihari@buffalo.edu

- Recurrent Neural Networks
 1. Unfolding Computational Graphs
 2. Recurrent Neural Networks
 3. Bidirectional RNNs
 4. Encoder-Decoder Sequence-to-Sequence Architectures
 5. Deep Recurrent Networks
 6. Recursive Neural Networks
 7. The Challenge of Long-Term Dependencies
 8. Echo-State Networks
 9. Leaky Units and Other Strategies for Multiple Time Scales
 10. LSTM and Other Gated RNNs
 11. Optimization for Long-Term Dependencies
 12. Explicit Memory

RNNs process sequential data


- Recurrent Neural Networks are a family of neural networks for processing sequential data
- RNN and CNN are both specialized architectures
- Just as CNN is specialized for processing grid of values, e.g., image
 - RNN is specialized for processing a sequence of values $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(\tau)}$
- Just as CNNs can readily scale images with large width/height and process variable size images
 - RNNs can scale to much longer sequences than would be practical for networks without sequence-based specialization
 - RNNs can also process variable-length sequences

Examples of Sequential Data and Tasks

- Sequence-to-sequence

- Speech recognition

- decompose sound waves into frequency and amplitude using Fourier transforms yielding a spectrogram shown




- Named Entity Recognition

- **Input:** Jim bought 300 shares of Acme Corp. in 2006
 - **NER:** [Jim]_{Person} bought 300 shares of [Acme Corp.]_{Organization} in [2006]_{Time}

- Sequence-to-symbol

- Sentiment
 - Speaker recognition

Neural network for 1-D convolution



Kernel $g(t)$:

$[\dots 0, w_1, w_0, 0\dots]$

Equations for outputs of this network:

$$y_0 = \sigma(W_0x_0 + W_1x_1 - b)$$

$$y_1 = \sigma(W_0x_1 + W_1x_2 - b) \text{ etc. upto } y_8$$

Note that kernel gets flipped in convolution

We can also write the equations in terms of elements of a general 8×8 weight matrix W as:


$$y_0 = \sigma(W_{0,0}x_0 + W_{0,1}x_1 + W_{0,2}x_2\dots)$$

$$y_1 = \sigma(W_{1,0}x_0 + W_{1,1}x_1 + W_{1,2}x_2\dots)$$

where $W = \begin{bmatrix} w_0 & w_1 & 0 & 0 & \dots \\ 0 & w_0 & w_1 & 0 & \dots \\ 0 & 0 & w_0 & w_1 & \dots \\ 0 & 0 & 0 & w_0 & \dots \\ \dots & \dots & \dots & \dots & \dots \end{bmatrix}$

Time Delay Neural Networks


- Time-delay neural networks perform convolution across 1-D temporal sequence
 - Convolution operation allows a network to share parameters across time, but is shallow
 - Each member of output is dependent upon a small no. of neighboring members of the input
 - Parameter sharing manifests in the application of the same convolutional kernel at each time step



A TDNN remembers the previous few training examples and uses them as input into the network. The network then works like a feed-forward, back propagation network.


RNN vs. TDNN

- RNNs share parameters in a different way
 - Each member of output is a function of previous members of output
 - Each output produced using same update rule applied to previous outputs
 - This recurrent formulation results in sharing of parameters through a very deep computational graph
- An unrolled RNN



RNN as a network with cycles

- An RNN is a class of neural networks where connections between units form a directed cycle
- This creates an internal state of the network which allows it to exhibit dynamic temporal behavior
- The internal memory can be used to process arbitrary sequences of inputs




Three layer network with input x , hidden layer z and output y
Context units c maintain a copy of the previous value of the hidden units

RNNs share same weights across Time Steps


- To go from multi-layer networks to RNNs:
 - Need to share parameters across different parts of a model
 - Separate parameters for each value of cannot generalize to sequence lengths not seen during training
 - Share statistical strength across different sequence lengths and across different positions in time
- Sharing important when information can occur at multiple positions in the sequence
 - Given “*I went to Nepal in 1999* ” and “*In 1999, I went to Nepal*”, an ML method to extract year, should extract 1999 whether in position 6 or 2
 - A feed-forward network that processes sentences of fixed length would have to learn all of the rules of language separately at each position
 - An RNN shares the same weights across several time steps

Problem of Long-Term Dependencies

- Easy to predict last word in “the clouds are in the *sky*,”
 - When gap between relevant information and place that it’s needed is small, RNNs can learn to use the past information



- “I grew up in France... I speak fluent *French*.”
 - We need the context of France, from further back.
 - Large gap between relevant information and point where it is needed



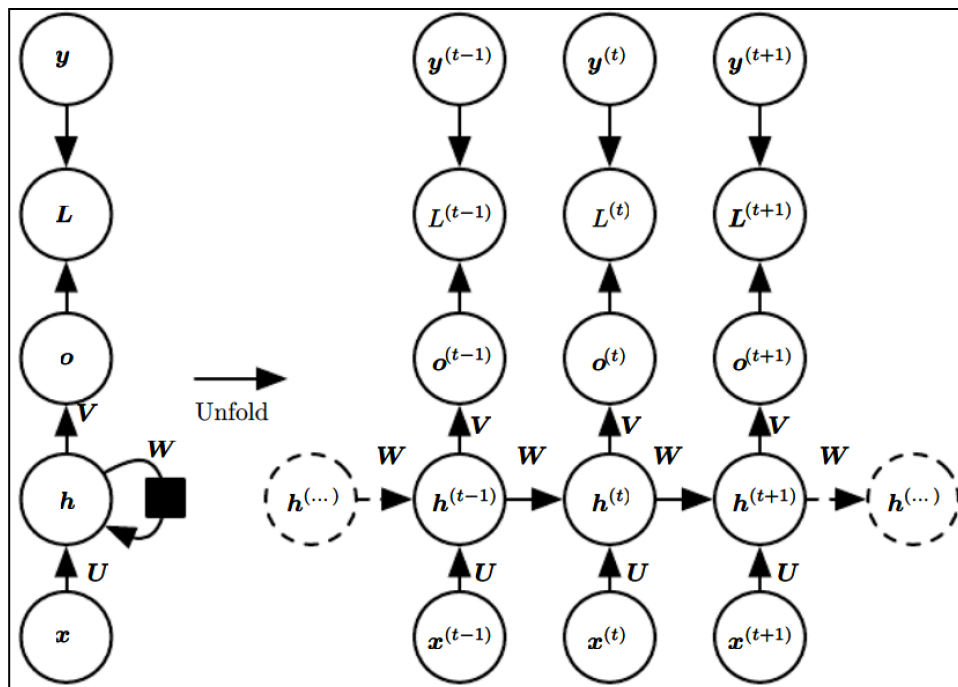
- In principle RNNs can handle it, but fail in practice
 - LSTMs offer a solution

RNN operating on a sequence

- RNNs operate on a sequence that contain vector $\mathbf{x}^{(t)}$ with time step index t , ranging from 1 to τ
 - Sequence: $\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(\tau)}$
 - RNNs operate on minibatches of sequences of length τ
- Some remarks about sequences
 - The steps need not refer to passage of time in the real world
 - RNNs can be applied in two-dimensions across spatial data such as image
 - Even when applied to time sequences, network may have connections going backwards in time, provided entire sequence is observed before it is provided to network

Computational Graphs for RNNs

- We extend computational graphs to include cycles
 - Cycles represent the influence of the present value of a variable on its own value at a future time step
 - In a Computational graph nodes are variables/operations
 - RNN to map input sequence of x values to output sequence of o values
 - Loss L measures how far each output o is from the training target y



- Forward propagation is given as follows:

For each time step t , $t=1$ to $t=\tau$

Apply the following equations


$$o^{(t)} = c + Vh^{(t)}$$

$$h^{(t)} = \tanh(a^{(t)})$$


$$a^{(t)} = b + Wh^{(t-1)} + Ux^{(t)}$$

Recurrent Neural Network

RNN



Unrolled RNN



Definition



inputs : $x = (x_1, x_2, \dots, x_T), x_i \in \mathbb{R}^I$
 hidden units : $h = (h_1, h_2, \dots, h_T), h_i \in \mathbb{R}^J$
 outputs : $y = (y_1, y_2, \dots, y_T), y_i \in \mathbb{R}^K$
 nonlinearity : \mathcal{H}

Activation Functions

$$h_t = \mathcal{H}(W_{xh}x_t + W_{hh}h_{t-1} + b_h)$$

$$y_t = W_{hy}h_t + b_y$$

Bidirectional RNN




inputs : $x = (x_1, x_2, \dots, x_T), x_i \in \mathbb{R}^I$
 hidden units : \vec{h} and \overleftarrow{h}
 outputs : $y = (y_1, y_2, \dots, y_T), y_i \in \mathbb{R}^K$
 nonlinearity : \mathcal{H}

$$\vec{h}_t = \mathcal{H}(W_{x\vec{h}}x_t + W_{\vec{h}\vec{h}}\vec{h}_{t-1} + b_{\vec{h}})$$



$$\overleftarrow{h}_t = \mathcal{H}(W_{x\overleftarrow{h}}x_t + W_{\overleftarrow{h}\overleftarrow{h}}\overleftarrow{h}_{t-1} + b_{\overleftarrow{h}})$$

$$y_t = W_{\overleftarrow{h}y}\overleftarrow{h}_t + W_{\vec{h}y}\vec{h}_t + b_y$$

Deep Bidirectional RNN




LSTM




- In deep learning:
 - Tasks of interest:
 - Classification
 - Feature learning
 - Method of learning
 - Backpropagation and gradient descent
- In graphical models:
 - Tasks of interest:
 - Transfer learning
 - Latent variable inference
 - Methods of learning
 - Parameter learning methods
 - Structure learning methods
- Hybrid graphical models combine the two types of models
 - They are trained using backpropagation

Hybrid Graphical Models and Neural Networks

- Hybrid NN and HMM



- Hybrid RNN+HMM



- Hybrid CNN+CRF

